

Lauren Pothuru

+1 (603) 714-1351 | laurenpothuru100@gmail.com

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SKILLS

Languages: Python, TypeScript, Rust, SQL, C, C++, Java

Tools/Libraries: React/React Native, AWS, BigQuery, PostgreSQL, LangChain, Inngest, PyTorch

EDUCATION

Cornell University *Ithaca, NY*

Aug 2024 – May 2028

B.A. in Computer Science • GPA: 3.6 / 4.0

Relevant Coursework: Machine Learning, Algorithms, Operating Systems, Computer System Organization, Data Structures

EXPERIENCE

Software Engineer Intern, Junior *New York, NY*

June 2026 – Aug 2026

- Designed and deployed an IAM-protected online-evaluation service with a BigQuery results layer and React UI, sampling 1% of a ~3,000-call/day platform; caught a recall regression on a new feature before its full rollout.
- Shipped desktop features to all users across React Native and Rust, including call trimming, live diarization and speaker naming, drag-and-drop interview guides, and Juni, an in-call assistant with live-transcript suggestions.
- Built a hybrid transcription pipeline combining commercial speech-to-text with LLM semantic correction and entity boosting, improving accuracy 5–18% across WER, numbers, entities, and KTAs; rolled it out to users and used online evals to flag hallucinations.

AI Research Fellow, Anote *New York, NY*

June 2026 – Aug 2026

- Built AgenticRAG-FP, an interventional benchmark for causal failure attribution in agentic RAG: inject a certified fault at a retrieval hop, re-execute the trajectory, and score diagnosers against the known cause; paper under review at EMNLP 2026 REALM.
- Built the benchmark harness and LLM evaluation platform, with judge-free answer-outcome labeling from certified span pairs.

Software Engineer Intern, InvestMint Inc. *Toronto, ON*

Dec 2025 – Present

- Built AWS-based ML infrastructure ingesting 10+ financial data sources for cash-flow forecasting for small and mid-sized businesses; cut infrastructure costs ~\$1K/month through a targeted vendor migration.

Teaching Assistant, Cornell Bowers CIS *Ithaca, NY*

Jun 2025 – Present

- Support 260+ students in algorithm design, runtime analysis, and correctness proofs; co-lead weekly sections and develop and grade problem sets.

RESEARCH & PROJECTS

“Retrieval Perturbation as a Stress-Test for RAG Faithfulness”, arXiv preprint

March 2026

- Designed a counterfactual benchmark to stress-test LLM failures, running 2,400 evaluations across four open-weight models with statistical significance testing.
- Identified over-abstention (not hallucination) as the dominant failure mode and released a three-archetype failure taxonomy and open-source benchmark.

AI Credit Analysis Agent, QuickFi *Fairport, NY*

Jan 2026 – May 2026

- Built a LangChain agent for a commercial equipment finance lender that validates credit applications against pulled records, calculates financial ratios, generates risk assessments, and flags fraudulent documents using LLM reasoning over normalized data.